

Stanford CS224W: GNNs for Recommender Systems

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



Stanford CS224W: Recommender Systems: Task and Evaluation

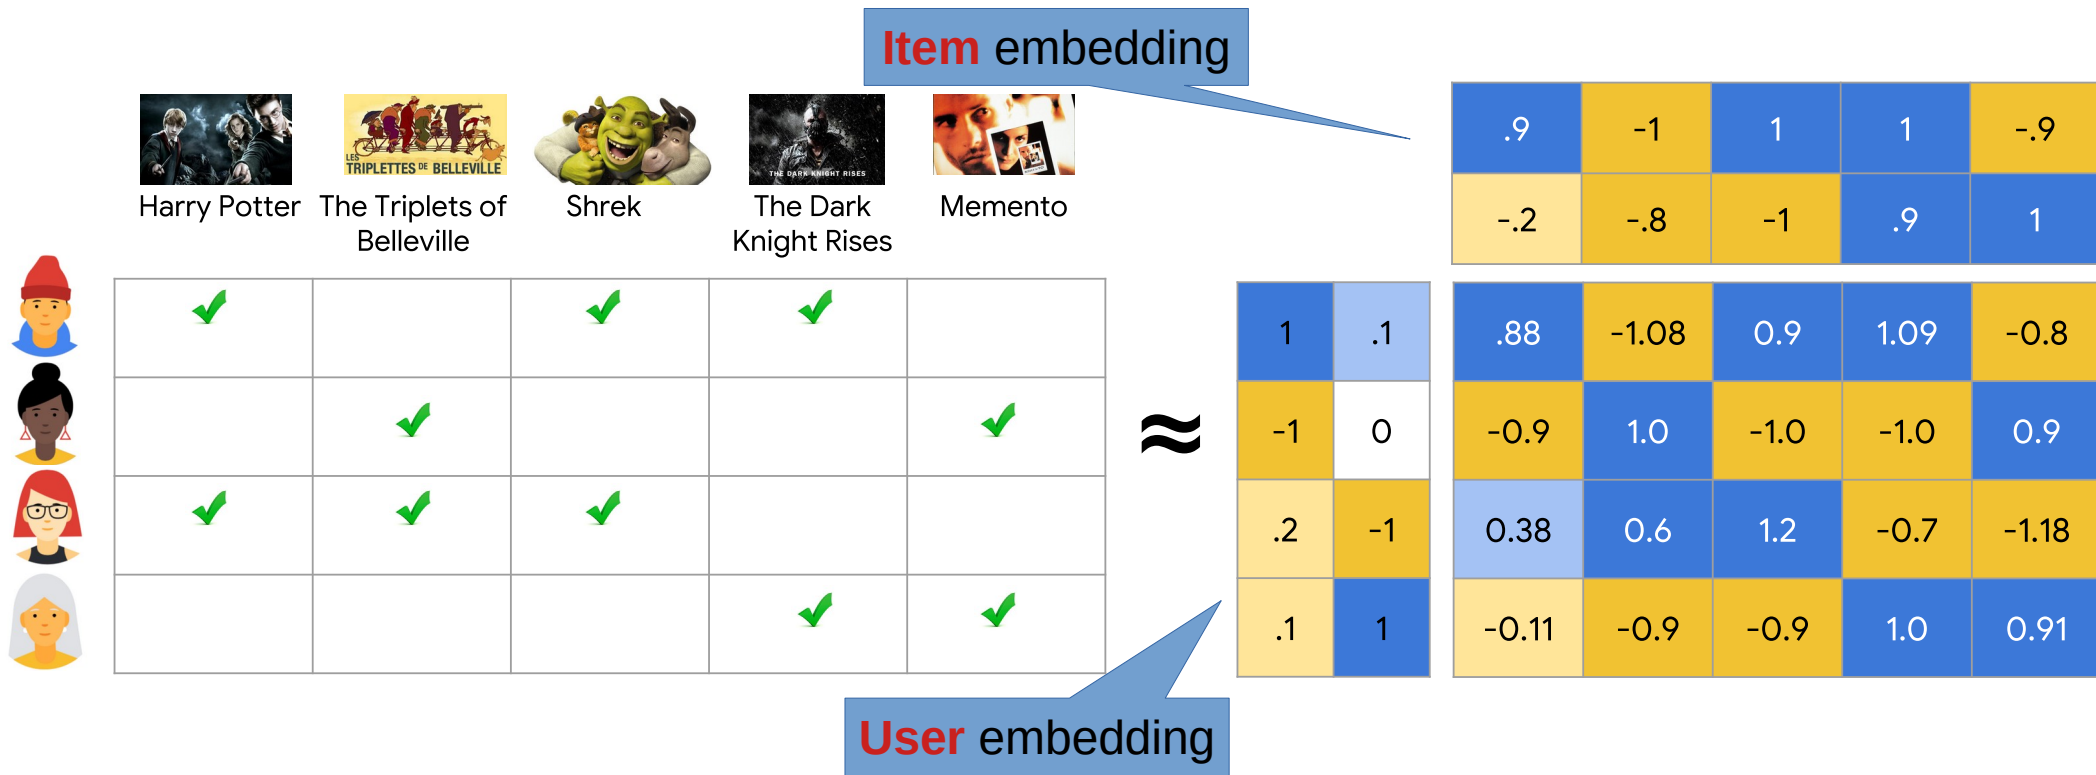
CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



Preliminary of Recommendation

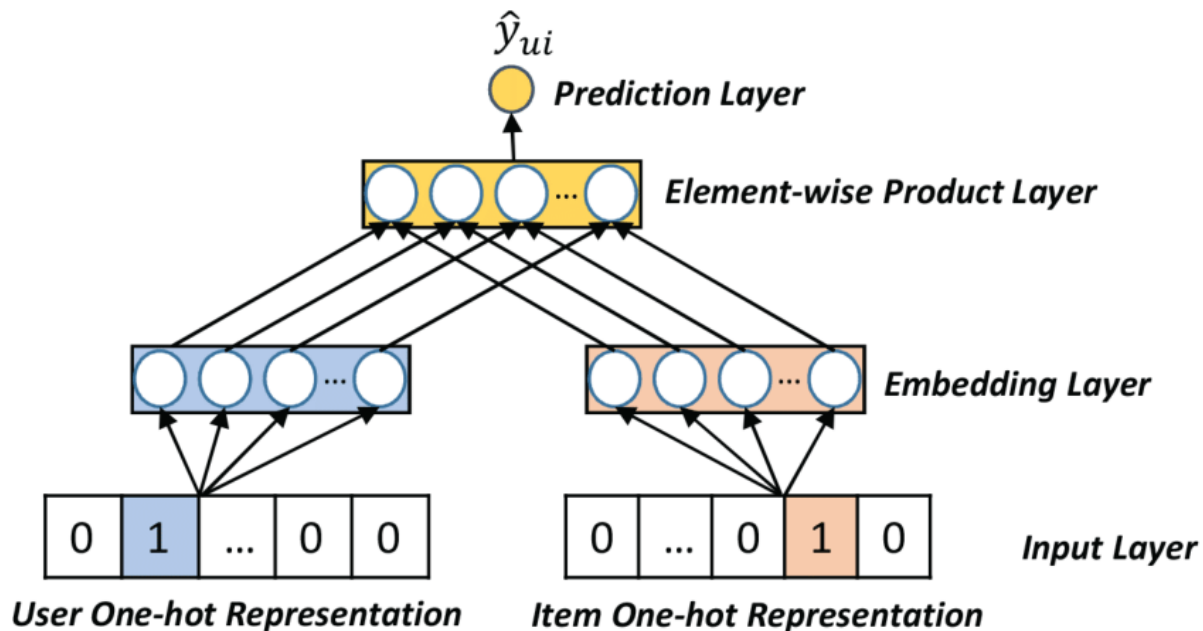
- **Information Explosion in the era of Internet**
 - 10K+ movies in Netflix
 - 12M products in Amazon
 - 70M+ music tracks in Spotify
 - 10B+ videos on YouTube
 - 200B+ pins (images) in Pinterest
- **Personalized recommendation (i.e., suggesting a small number of interesting items for each user)** is critical for users to effectively explore the content of their interest.

Matrix Factorization




The embeddings are learned such that their dot product is a good approximation of the user-item matrix

Matrix-factorization as a shallow neural network model



Recommender System as a Graph

- Recommender system can be naturally modeled as a **bipartite graph**
 - A graph with two node types: **users** and **items**.
 - **Edges** connect users and items
 - Indicates user-item interaction (e.g., click, purchase, review etc.)
 - Often associated with timestamp (timing of the interaction).




Recommendation Task

■ Given

- Past user-item interactions

■ Task

- Predict new items each user will interact in the future.
- Can be cast as **link prediction** problem.
 - Predict new user-item interaction edges given the past edges.



Stanford CS224W: Recommender Systems: Embedding-Based Models

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>




Notation

■ Notation:

- U : A set of all users
- V : A set of all items
- E : A set of observed user-item interactions
 - $E = \{(u, v) \mid u \in U, v \in V, u \text{ interacted with } v\}$

Score Function


- To get the top- K items, we need a score function for user-item interaction:
 - For $u \in U$, $v \in V$, we need to get a real-valued scalar $\text{score}(u, v)$.
 - **K items with the largest scores for a given user u** (excluding **already-interacted items**) are then recommended.



For $K = 2$, **recommended items** for user u would be $\{v_1, v_3\}$.

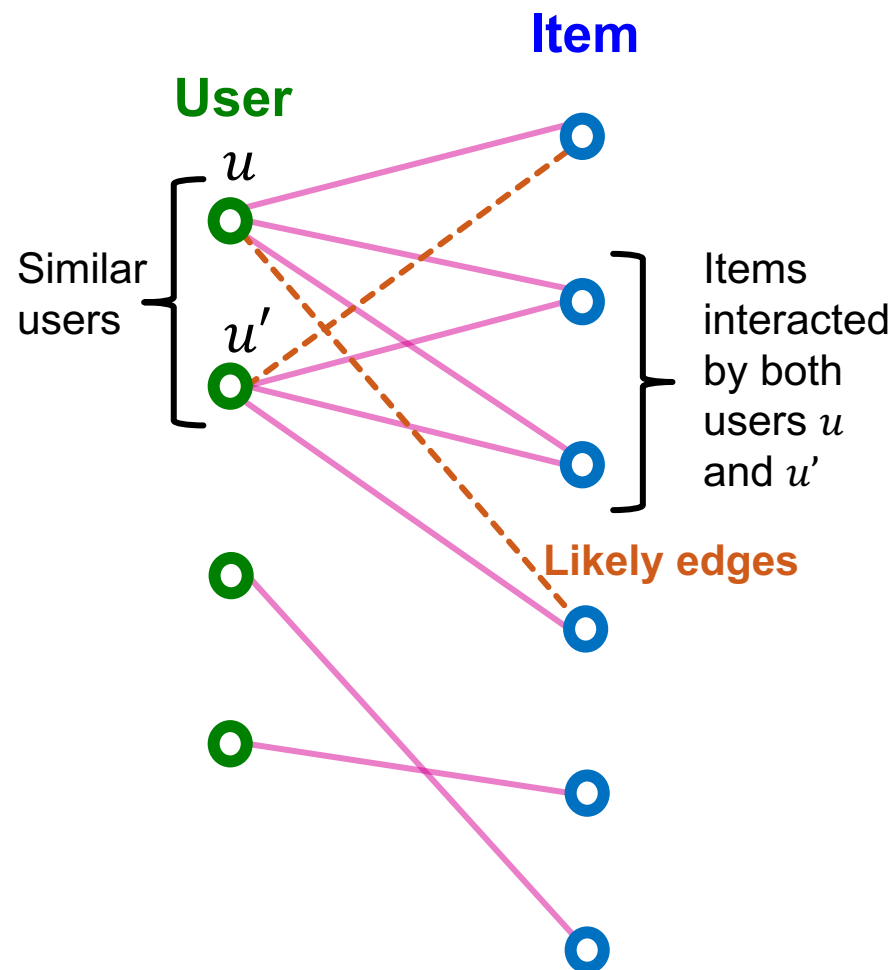
Embedding-Based Models

- We consider **embedding-based models** for scoring user-item interactions.
 - For each user $u \in U$, let $\mathbf{u} \in \mathbb{R}^D$ be its D -dimensional embedding.
 - For each item $v \in V$, let $\mathbf{v} \in \mathbb{R}^D$ be its D -dimensional embedding.
 - Let $f_\theta(\cdot, \cdot): \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$ be a parametrized function.
 - Then, $\text{score}(u, v) \equiv f_\theta(\mathbf{u}, \mathbf{v})$



Why Embedding Models Work?

- **Underlying idea:**
Collaborative filtering
 - Recommend items for a user by **collecting preferences of many other similar users.**
 - **Similar users tend to prefer similar items.**
- **Key question:** **How to capture similarity between users/items?**



Why Embedding Models Work?

- Embedding-based models can capture similarity of users/items!
 - **Low-dimensional embeddings *cannot* simply memorize all user-item interaction data.**
 - Embeddings are forced to **capture similarity between users/items to fit the data.**
 - This allows the models to make effective prediction on *unseen* user-item interactions.

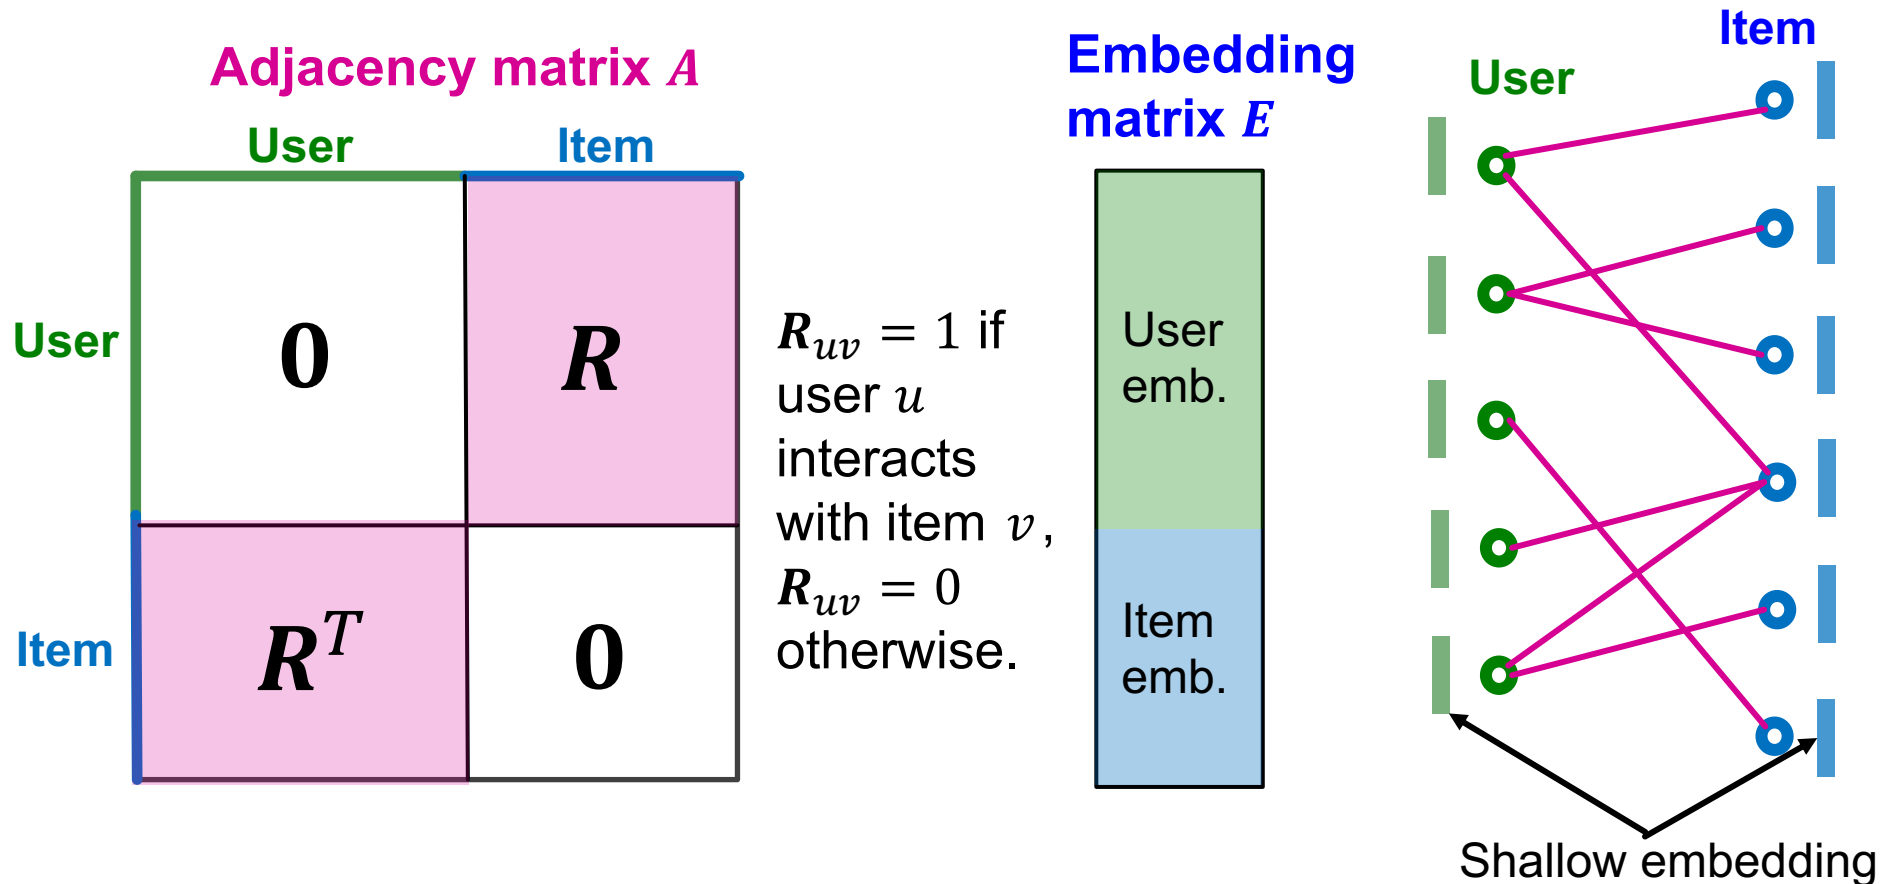
Stanford CS224W: LightGCN

CS224W: Machine Learning with Graphs
Jure Leskovec, Stanford University
<http://cs224w.stanford.edu>



Adjacency and Embedding Matrices

- **Adjacency matrix** of a (undirected) bipartite graph.
- **Shallow embedding matrix**.



Matrix Formulation of GCN

- **Recall:** Diffusion matrix of C&S.
- Let \mathbf{D} be the degree matrix of \mathbf{A} .
- Define the normalized adjacency matrix $\tilde{\mathbf{A}}$ as

$$\tilde{\mathbf{A}} \equiv \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$$

Note: Different from the original GCN, self-connection is omitted here.


- Let $\mathbf{E}^{(k)}$ be the embedding matrix at k -th layer.
- Each layer of GCN's aggregation can be written in a matrix form:

$$\mathbf{E}^{(k+1)} = \text{ReLU}(\underbrace{\tilde{\mathbf{A}} \mathbf{E}^{(k)}}_{\text{Neighbor aggregation}} \underbrace{\mathbf{W}^{(k)}}_{\text{Learnable linear transformation}})$$

Neighbor aggregation

Learnable linear transformation

Matrix of node embeddings $\mathbf{E}^{(k)}$



Each row stores node embedding

Simplifying GCN (1)

- Simplify GCN by **removing ReLU non-linearity**:

$$E^{(k+1)} = \tilde{A} E^{(k)} W^{(k)} \quad \text{Original idea from SGC [Wu et al. 2019]}$$

- The final node embedding matrix is given as

$$\begin{aligned}
 E^{(K)} &= \tilde{A} \underbrace{E^{(K-1)}}_{\text{pink}} W^{(K-1)} \\
 &= \tilde{A} \left(\underbrace{\tilde{A} E^{(K-2)}}_{\text{green}} W^{(K-2)} \right) W^{(K-1)} \\
 &= \tilde{A} \left(\tilde{A} \left(\dots \left(\underbrace{\tilde{A} E^{(0)}}_{\text{green}} W^{(0)} \right) \dots \right) W^{(K-2)} \right) W^{(K-1)} \\
 &= \tilde{A}^K \underbrace{E}_{\text{blue}} \left(\underbrace{W^{(0)} \dots W^{(K-1)}}_{\text{orange}} \right)
 \end{aligned}$$

Set E as input embedding $E^{(0)}$

Simplifying GCN (2)

- Removing ReLU significantly simplifies GCN!

$$\mathbf{E}^{(K)} = \boxed{\tilde{\mathbf{A}}^K \mathbf{E}} \mathbf{W}$$

Diffusing node embeddings
along the graph

(similar to C&S that diffuses
soft labels along the graph)

$$\mathbf{W} \equiv \mathbf{W}^{(0)} \dots \mathbf{W}^{(K-1)}$$

- **Algorithm:** Apply $\mathbf{E} \leftarrow \tilde{\mathbf{A}} \mathbf{E}$ for K times.
 - Each matrix multiplication diffuses the current embeddings to their one-hop neighbors.
 - **Note:** $\tilde{\mathbf{A}}^K$ is dense and never gets materialized. Instead, the above iterative matrix-vector product is used to compute $\tilde{\mathbf{A}}^K \mathbf{E}$.

Multi-Scale Diffusion

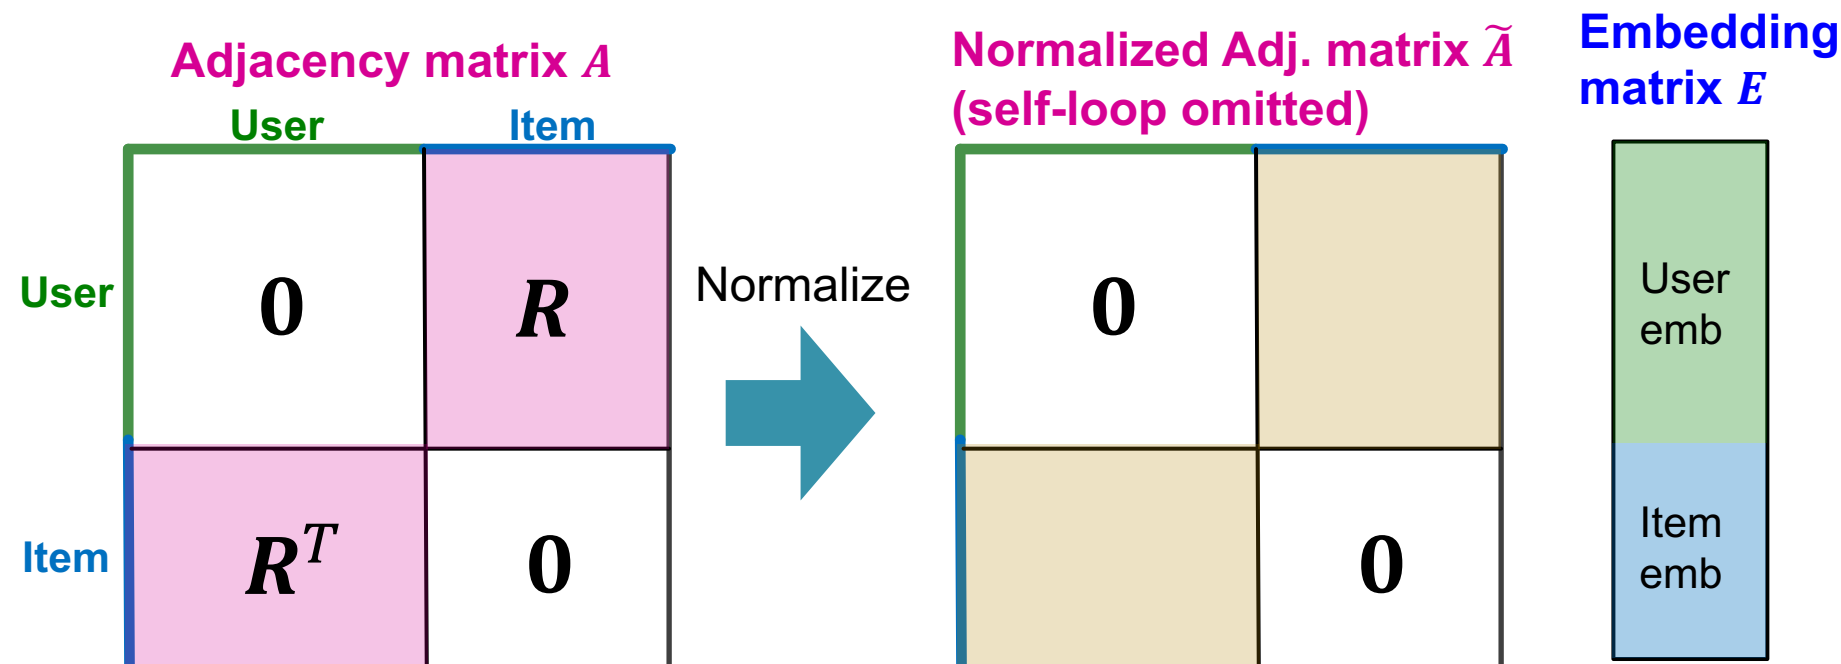
- We can consider **multi-scale diffusion**

$$\alpha_0 E^{(0)} + \alpha_1 E^{(1)} + \alpha_2 E^{(2)} + \dots + \alpha_K E^{(K)}$$

- The above includes embeddings diffused at multiple hop scales.
- $\alpha_0 E^{(0)} = \alpha_0 \tilde{\mathbf{A}}^0 E^{(0)}$ acts as a self-connection (that is omitted in the definition $\tilde{\mathbf{A}}$)
- The coefficients, $\alpha_0, \dots, \alpha_K$, are hyper-parameters.
- For simplicity, LightGCN uses the uniform coefficient, i.e., $\alpha_k = \frac{1}{K+1}$ for $k = 0, \dots, K$.

LightGCN: Model Overview (1)

- **Given:**
 - Adjacency matrix A
 - Initial learnable embedding matrix E




LightGCN: Model Overview (2)


- Iteratively diffuse embedding matrix E using \tilde{A}

For $k = 0 \dots K - 1$,


Embedding
matrix $E^{(k+1)}$



Normalized Adj. matrix \tilde{A}
(self-loop omitted)



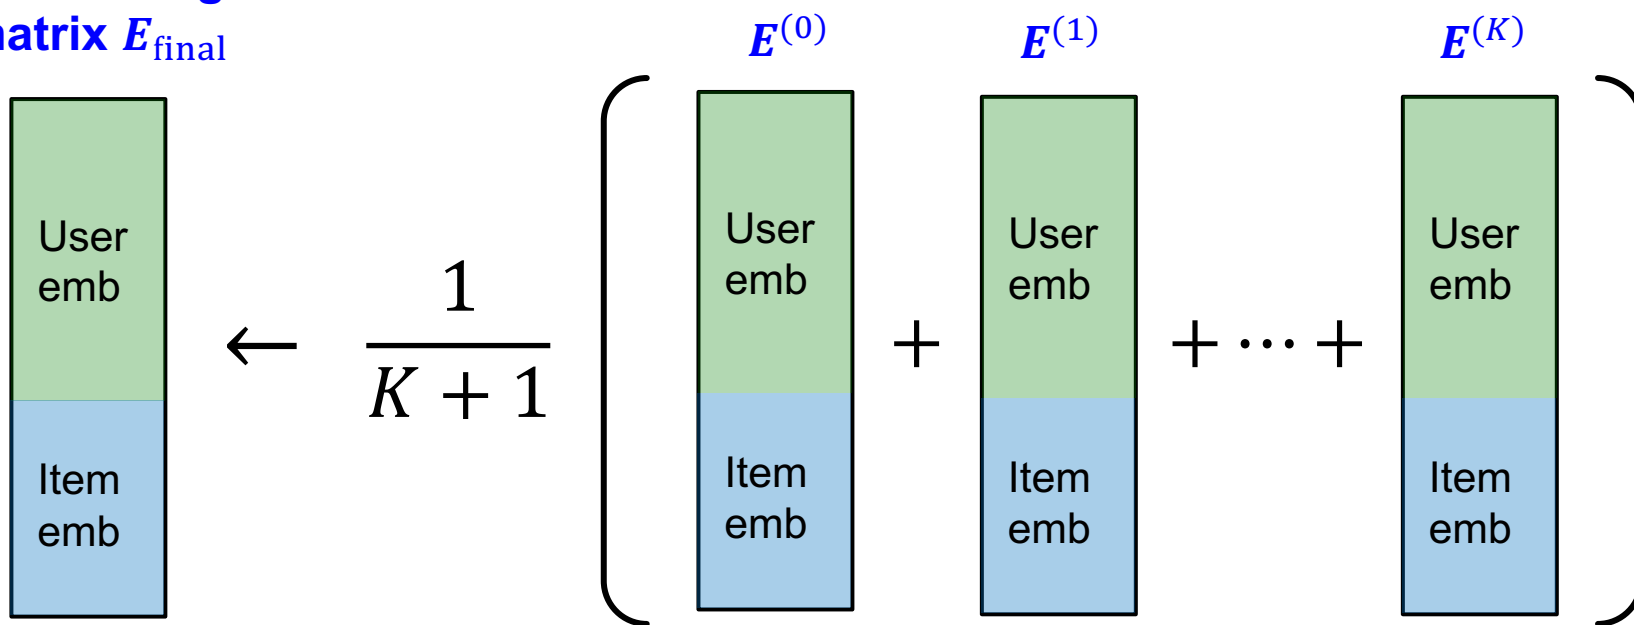
Embedding
matrix $E^{(k)}$
($E^{(0)}$ is set to E)



LightGCN: Model Overview (3)

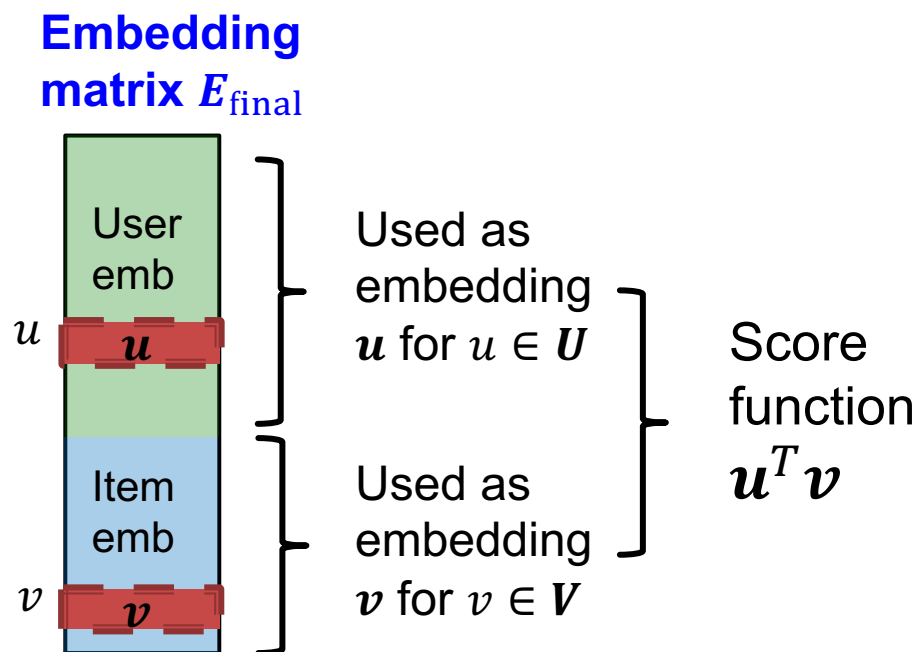
- Average the embedding matrices at different scales.

Embedding
matrix E_{final}



LightGCN: Model Overview (4)

- **Score function:**
 - Use user/item vectors from E_{final} to score user-item interaction



LightGCN: Intuition

- **Question:** Why does the simple diffusion propagation work well?
- **Answer:** The diffusion directly encourages the embeddings of similar users/items to be similar.
 - Similar users share many common neighbors (items) and are expected to have similar future preferences (interact with similar items).

LightGCN and GCN/C&S

- The embedding propagation of LightGCN is closely related to GCN/C&S.
- **Recall:** GCN/C&S (neighbor aggregation part)

$$\mathbf{h}_v^{(k+1)} = \sum_{u \in N(v)} \frac{1}{\sqrt{d_u} \sqrt{d_v}} \cdot \mathbf{h}_u^{(k)}$$

Node degree

- Self-loop is added in the neighborhood definition.
- LightGCN uses the same equation except that
 - Self-loop is *not* added in the neighborhood definition.
 - Final embedding takes the average of embeddings from all the layers: $\mathbf{h}_v = \frac{1}{K+1} \sum_{k=0}^K \mathbf{h}_v^{(k)}$.

LightGCN and MF: Comparison

- Both LightGCN and Matrix Factorization (MF) **learn a unique embedding for each user/item.**
- The difference is that
 - MF directly uses the shallow user/item embeddings for scoring.
 - LightGCN uses the *diffused* user/item embeddings for scoring.
- LightGCN performs better than MF but are also more computationally expensive due to the additional diffusion step.
 - The final embedding of a user/item is obtained by aggregating embeddings of its multi-hop neighboring nodes.

LightGCN: Summary

- LightGCN simplifies NGCF by **removing the learnable parameters of GNNs.**
- **Learnable parameters are all in the shallow input node embeddings.**
 - Diffusion propagation only involves matrix-vector multiplication.
 - The simplification leads to better empirical performance than NGCF.